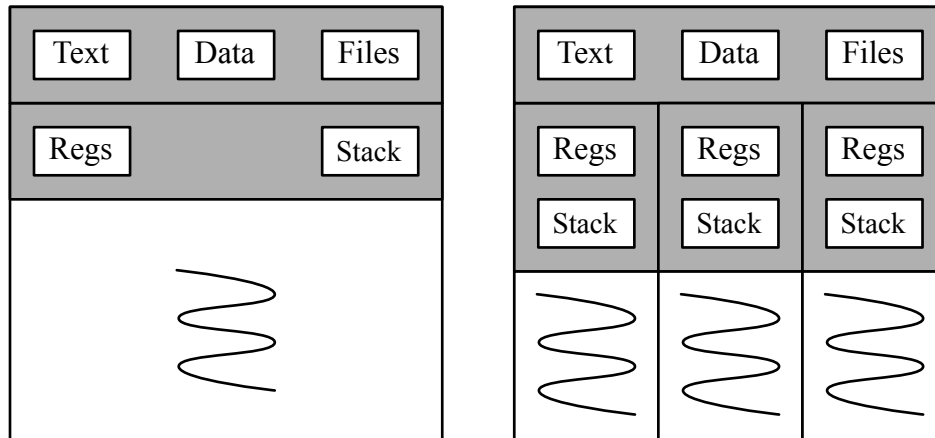


Lecture Notes 4

Threads

- Threading Models
 - Single Threaded – Single sequential execution context
 - Multithreaded – Multiple concurrent execution contexts



- Multithreaded Benefits
 - Responsiveness – Allows interactive applications to respond when other thread is blocked
 - Resource Sharing – Files and memory are automatically shared in multithreaded programs. Multi-process requires some IPC mechanism.
 - Economy – Lower resource cost, possibly lower overhead of context switch.
 - Scalability – Heavy weight threads may take advantage of multiple cores, single threaded process can't take advantage of multiple cores.
- Multicore Programming (or Multiprocessor Programming)
 - Parallelism – Multiple Tasks (or Threads) run simultaneously
 - Concurrency – Multiple Tasks (or Threads) continue to make progress
 - Amdahl's Law – Formula for potential speedup of multiple processing cores.
 - Programming Challenges
 - Identifying Tasks – Examine application to divide up concurrent tasks
 - Balance – Identifying and splitting up work equally
 - Data Splitting – Data accessed must be divided up between separate cores
 - Data Dependency – Must ensure dependencies are synchronized
 - Testing and Debugging – Increased difficulty due non-determinism
 - Types of Parallelism
 - Data Parallelism – Distributes subsets of data across multiple computing cores
 - Task Parallelism – Distributes data and tasks (threads) across multiple computing cores

- Multithreading Models
 - Heavy Weight Threads – Kernel space, visible to the OS, can concurrently executing on independent cores
 - Light Weight Threads – User space only, scheduled by application, all share a core
 - Fibers – Light Weight Threads that are cooperatively scheduled (non-preemptive threads)
 - Many-to-One Model – Many user space threads to one kernel space thread
 - One-to-One Model – One user space thread to one kernel space thread
 - Many-to-Many Model – Multiple user space threads to multiple (usually fewer) kernel space threads
 - Two-level Model – Many-to-Many that also allows user space to kernel space binding
- Thread Libraries – API for programmer to create and manage threads
 - Pthreads – POSIX standard library may be kernel or user space depending upon implementation
 - Windows Threads – Threading interface for Windows OS
 - Java Threads – Fundamental part of Java programming language
- Implicit Threading – Multiple threads are implied, not explicitly created
 - Thread Pools – Limited number of threads created and are taken from a pool
 - OpenMP – Open Multi-Processing library, compiler directives and API
 - Parallel Regions – A region of code that can execute in parallel
 - Grand Central Dispatch (GCD) – Apple technology similar to OpenMP
 - Blocks – A self contained unit of work
 - Dispatch Queue – The queue that holds the blocks
 - Main Queue – Per process serial queue
- Threading Issues
 - System Calls (fork() and exec())
 - fork() – Should all threads be duplicated?
 - exec() – If called after fork(), thread duplication not needed.
 - Signal Handling – Signals notify a process of an event is “handled” by a handler function
 - Default Signal Handler – Default handler defined by the kernel
 - User-Defined Signal Handler – User specified function to handle the signal
 - Which thread handles the signal?
 - What if all need to receive the signal?
 - Thread Cancellation – Termination of threads
 - Target Thread – Thread that is set to be cancelled
 - Asynchronous Cancellation – Thread immediately terminates Target Thread
 - Deferred Cancellation – Target Thread periodically checks for termination
 - Cancellation Point – Place where deferred cancellation occurs
 - Cleanup Handler – Function that is called to release of resources

- Thread-Local Storage – Thread own copy of data needed for processing
- Scheduler Activations – Communication point between user-thread library and kernel
 - Light Weight Process (LWP) – Each LWP is connected to a kernel thread
 - Upcall – Kernel notifying application about events
 - Upcall Handler – Function that handles the upcall event